



gen\_paxos

kuenishi

2009/7/3 3<sup>rd</sup> Erlang-meeting in  
Shirogane-takanawa, Tokyo

# 自己紹介

- ウエニシ コウタ(kuenishi@gmail.com)
- 趣味
  - 役に立つかどうか分かりもしないプログラムをつくること
  - ErlangやPython
  - See `id:kuenishi`, `g:erlang:id:kuenishi`, `@kuenishi`
    - <http://github.com/kuenishi>
    - <http://bitbucket.org/kuenishi>
- Recent Activities
  - Yet another TC-Erlang binding
    - <http://bitbucket.org/kuenishi/yatce/>
  - Mercurial l10n
    - 1.3 has been released on 7/1!
    - <http://bitbucket.org/foozy/mercurial-translation-ja/>

# Outline

- Introduction
  - Why PAXOS?
  - What is PAXOS?
  - How can we implement PAXOS?
    - Why Erlang?
  - Where is PAXOS?
  - When?
  - Appendix
- 
- Notice: Basically English, Sometimes Japanese

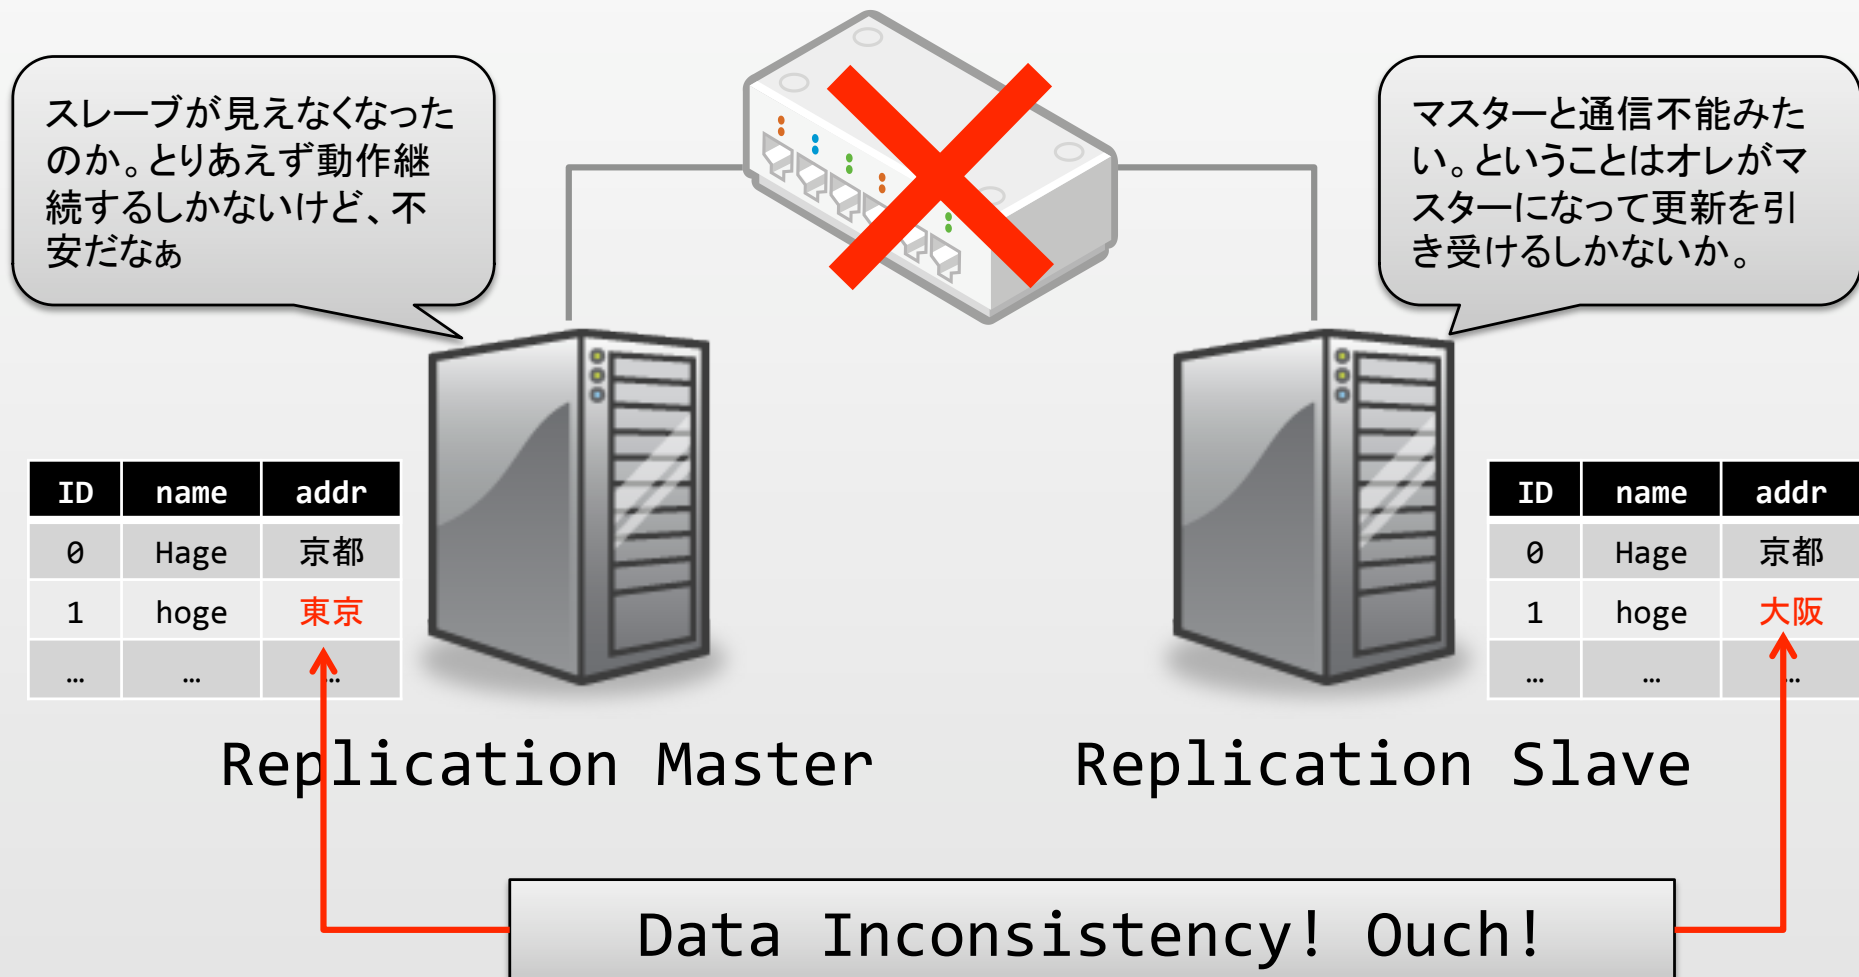
# Introduction – what is gen\_paxos?

- **GENeric PAXOS module written in Erlang**
  - Reduction of the cost of re-invention of wheels.
  - At first, I had aimed to make it sth like ‘behaviour’
  - [http://github.com/kuenishi/gen\\_paxos](http://github.com/kuenishi/gen_paxos)
- **Related works (I think not enough)**
  - libpaxos
  - gen\_leader
- **Why not enough?**

# Why PAXOS?

- Split Brain

- Typically, broken switches / wires
- Both master and slave updates independently



## Why PAXOS? – cont'd

- Split Brain Problem
  - Separate actors can't coordinate
    - Lacking coordination -> Multi-master
    - Multi-master -> Data Inconsistency
    - Inconsistency -> Tragedy
  - Typically, BigTable solely depends upon Chubby[2]
    - BigTable's Master Election, so on.
- Byzantine Generals Problem
  - A faulty actor who lies destroys consistency
  - I don't know current implementation of gen\_paxos solves this problem – not in scope
  - See Byzantine PAXOS in Wikipedia[6]

# What is PAXOS?

- **Distributed Consensus/Coordination Algorithm**
  - Masterless (no SPOF)
  - Robust to split brain environment such as network separation
  - Proved to converge within infinite time period[9]
  - Works iff majority of fixed group of actors is alive and communicable

## What is PAXOS? – Example

- Two phase election – phase 1
  - Larger  $n$  is prior



21:36  
新宿！

$n=3$



21:32  
今度の勉強会、  
どこでやる？



21:35       $n=2$   
横浜だろJK

21:36  
新宿ですかorz



$n=1$

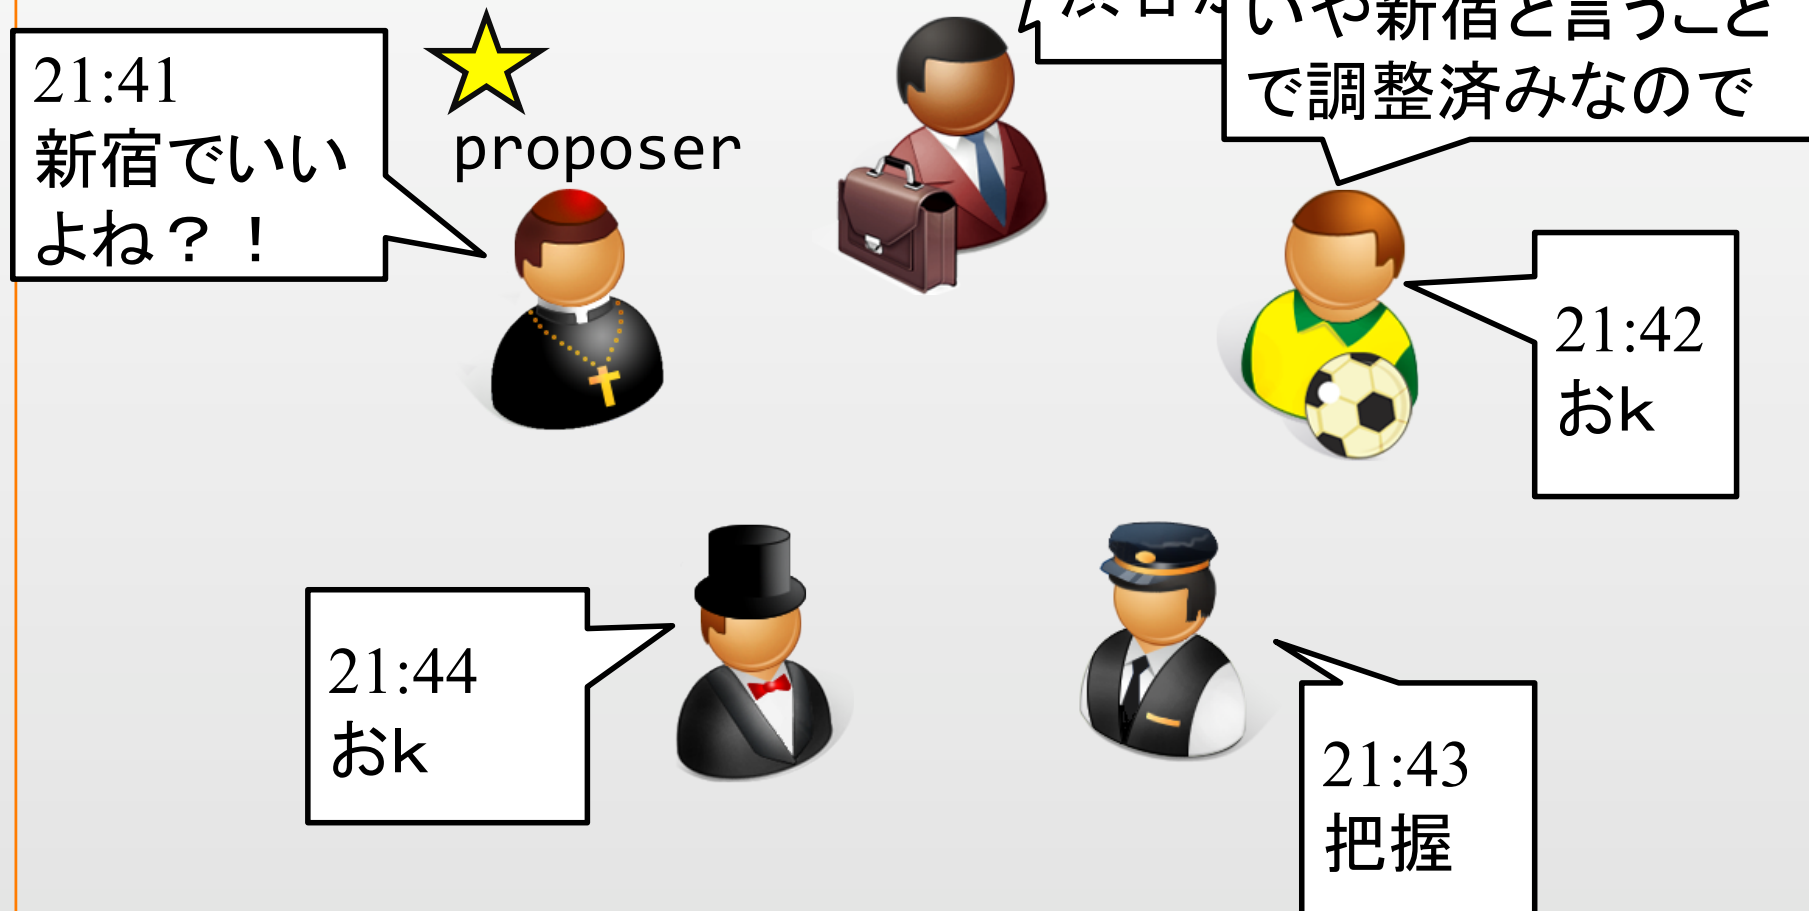
21:38  
新宿ですねorz



# What is PAXOS? – Example

- Two phase election – phase 2

- confirmation



# What is PAXOS? – data model in gen\_paxos

- **Actors**

- n processors with processor ID


- **Data**

- Instance ID (Subject)
- Proposal ID (N)
- Proposal itself (Value)

- N should be unique

- At first, V differs

- later V converges



ID	1	2	3	4	5
n(1)	1	2	3	4	5
n(2)	6	7	8	9	10
n(3)	11	12	13	14	15
n(4)	16	17	18	19	20
n(...)	...	...	..	...	..

# What is PAXOS? – details in gen\_paxos

- Basic sequence of Each actor
- Prepare Phase
  - If Prepare comes, agree or disagree (be Acceptor)
  - If Prepare doesn't come, send Prepare (be Proposer)
  - If majority of group agrees Prepare, send Propose
  - If timeout, increase n and restart Prepare phase
- Propose Phase
  - Acceptor – agree to Propose (ignoring lower n)
  - Proposer – decided if majority agrees
  - If timeout, increase n and restart Prepare phase
- Commit Phase
  - Send the committed Value again to notify

# What is PAXOS? – details in gen\_paxos

- Basic sequence of Each actor
- Prepare Phase
  - Prepareが来たら賛成したりしなかったり (Acceptor)
  - Prepareが来なかったらPrepareを送る (Proposer)
  - 過半数からPrepareの賛成が来たらPropose
  - タイムアウトしたらnを増やして戻る
- Propose Phase
  - Acceptor - Proposeが来たら賛成 (nの値によっては無視)
  - Proposer - 過半数から賛成が来たら決定
  - タイムアウトしたらnを増やしてPrepare Phaseに戻る
- Commit Phase
  - 決まった値をもう一度送る

# What is PAXOS? – details in gen\_paxos

- Basic sequence of Each actor

## • Pre

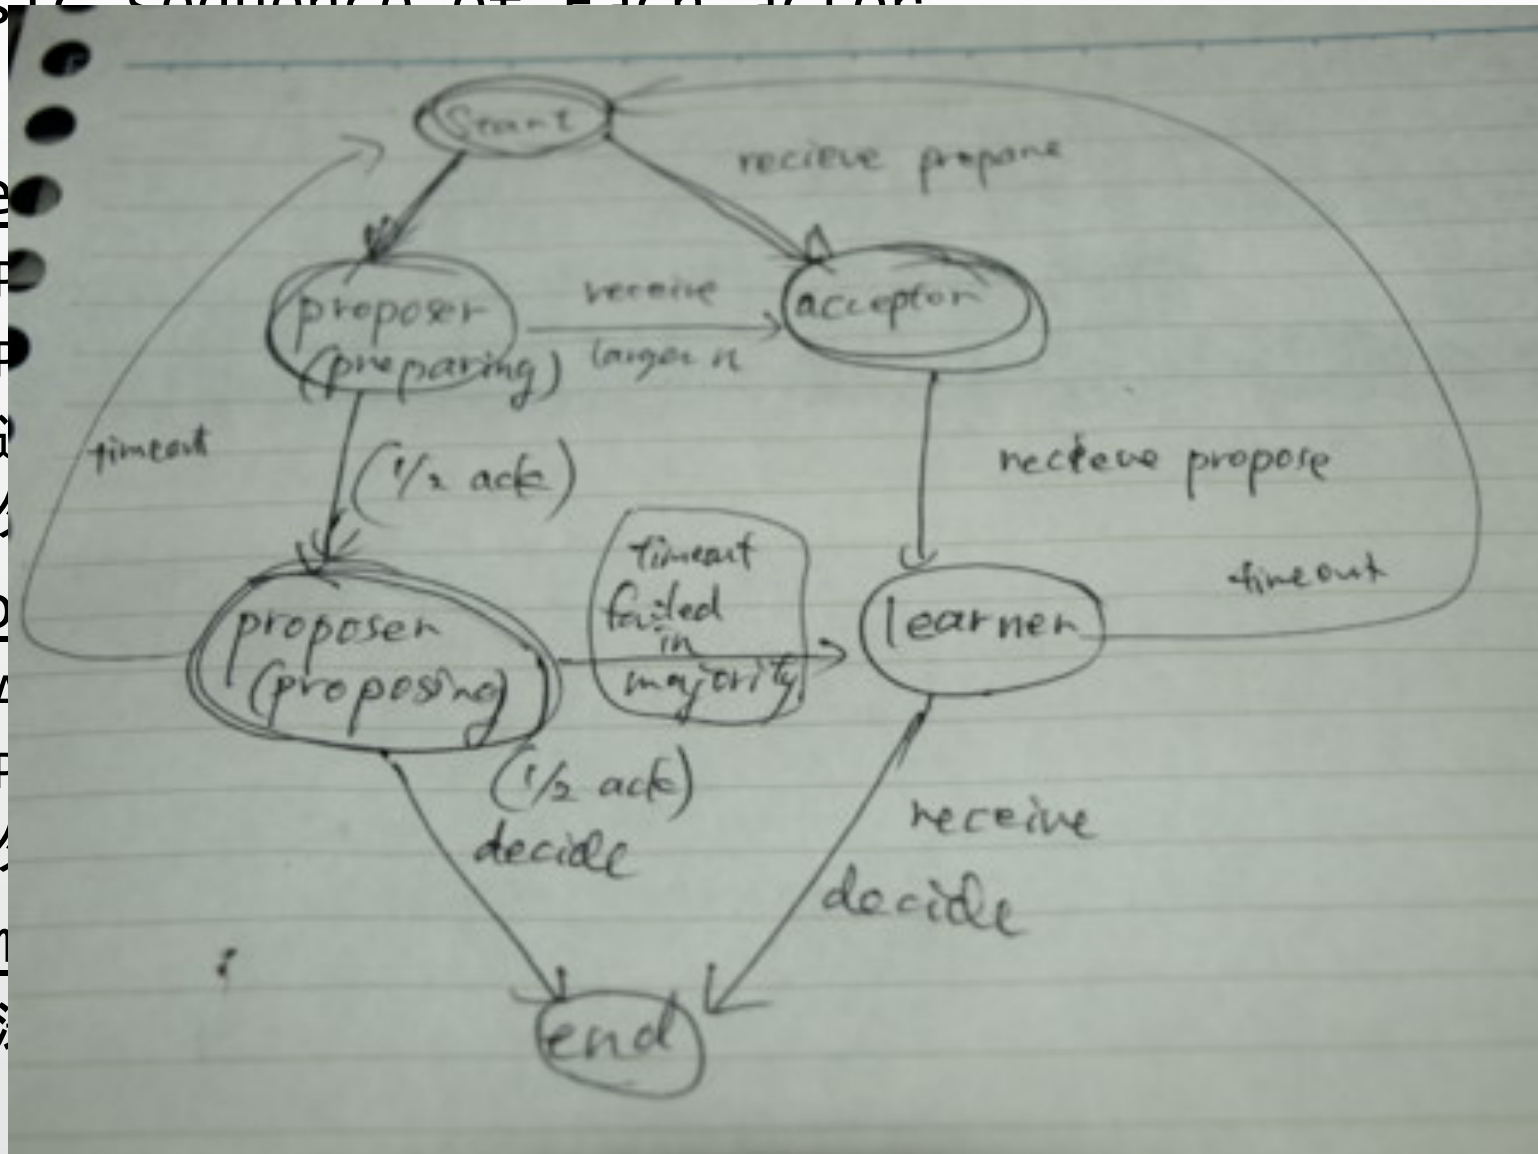
- P
- P
- 1
- 1

## • Pro

- A
- P
- 1

## • Com

- 1



How can we implement PAXOS?

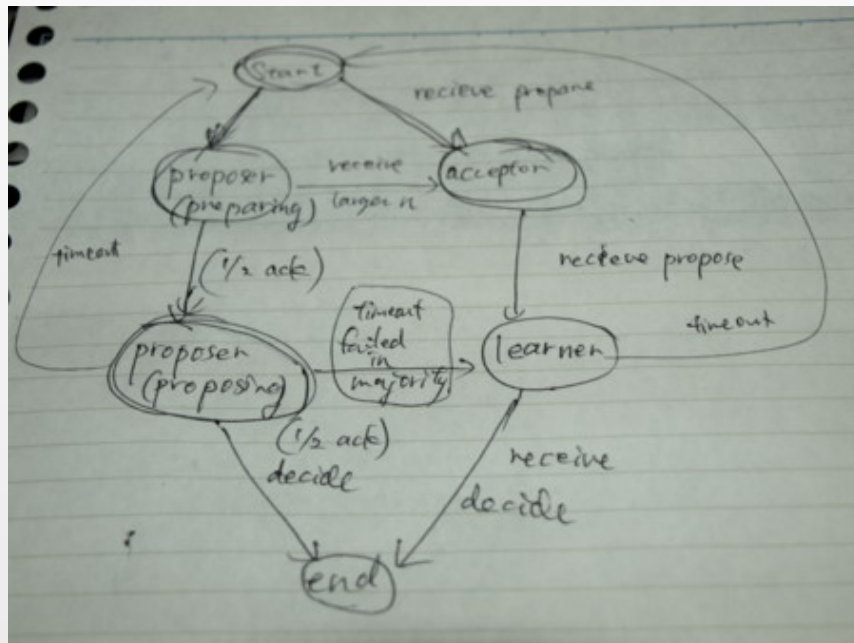
**1. Make a State Diagram**

**2. Write an FSM code**



# How can we implement PAXOS? (1/2)

## 1. Make a State Diagram



[State]

State	Event	Action	NextState
Start	receive propose	prepare(n, v)	proposer (preparing)
proposer (preparing)	receive larger n	ignore	proposer (preparing)
proposer (preparing)	timeout	prepare(n', v)	proposer (proposing)
proposer (preparing)	(1/2 ack)	prepare_result(n, v)	proposer (proposing)
proposer (proposing)	decide	end(n, v)	end
proposer (proposing)	decide	propose(n, v)	learner
learner	receive decide	end(n, v)	end
learner	timeout	prepare(n', v)	proposer (proposing)
learner	timeout failed in majority	prepare(n', v)	proposer (proposing)
learner	timeout	prepare(n', v)	proposer (proposing)

[Event]

{Action, NextState}

Define:

•{State, Event} -> {NextState, Action}

# How can we implement PAXOS? (2/2)

## 2. Write an FSM code



ignore?						
	A	B	C	D	E	F
1	ある議題についてのstate	states	通信	ブロードキャスト		
2	initial(n, v)	n' < n < v''	prepare(n, v)	acceptor(n, v)	learner(n, v)	decided(n, v)
3	prepare(n', v')	prepare(n, v)を送り出す	ignore	prepare_result(n, v)を送る	ignore	ignore
4	prepare(n', v')	prepare_result(0, nil)を送ってacceptor(n, v)へ	prepare_result(n, v)を送ってacceptor(n', v')へ	prepare_result(n, v)を送ってacceptor(n', v')へ	prepare_result(n, v)を送ってacceptor(n', v')へ	ignore
5	prepare_result(0, nil)	ignore	過半数を過ぎていたら propose(n, v)を送って propose(n, v)へ	ignore	ignore	ignore
6	prepare_result(n', v')	ignore	過半数を過ぎていたら propose(n, v)を送って propose(n, v)へ	ignore	ignore?	ignore
7	prepare_result(n, v)	ignore	過半数を過ぎていたら propose(n, v)を送って propose(n, v)へ	-	-	ignore
8	prepare_result(n', v')	ignore	prepare_result(n, v)を送ってacceptor(n', v')へ	prepare_result(n, v)を送ってacceptor(n', v')へ	ignore?	ignore
9	general_timeout	prepare(n, v)をみんなに送ってprepare(n, v)へ	n'', nilへ	n'', nilへ	n'', nilへ	n'', nilへ
10	propose(n', v')	ignore	ignore	ignore	ignore	decided(n, v)を送信
11	propose(n, v)	-	-	prepare_result(n, v)を送って learner(n, v)へ	ignore	ignore
12	propose(n'', v'')	ignore	prepare_result(n, v)を送って learner(n'', v'')へ	prepare_result(n, v)を送って learner(n'', v'')へ	prepare_result(n, v)を送って learner(n'', v'')へ	ignore
13	prepare_result(n', v')	ignore	ignore	ignore	ignore	reply decide(n, v)
14	prepare_result(n, v)	ignore	ignore	過半数を超えれば全員に decided(n, v)を送り end(n, v)へ	-	reply decide(n, v)
15	prepare_result(n'', v'')	ignore	prepare_result(n, v)を送って learner(n'', v'')へ	-	-	reply decide(n, v)
16	decide(n, v)	end(n, v)	end(n, v)	end(n, v)	end(n, v)	end(n, v)
17	decide(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')
18						
19						



Define a set of  
**Module:State/2**

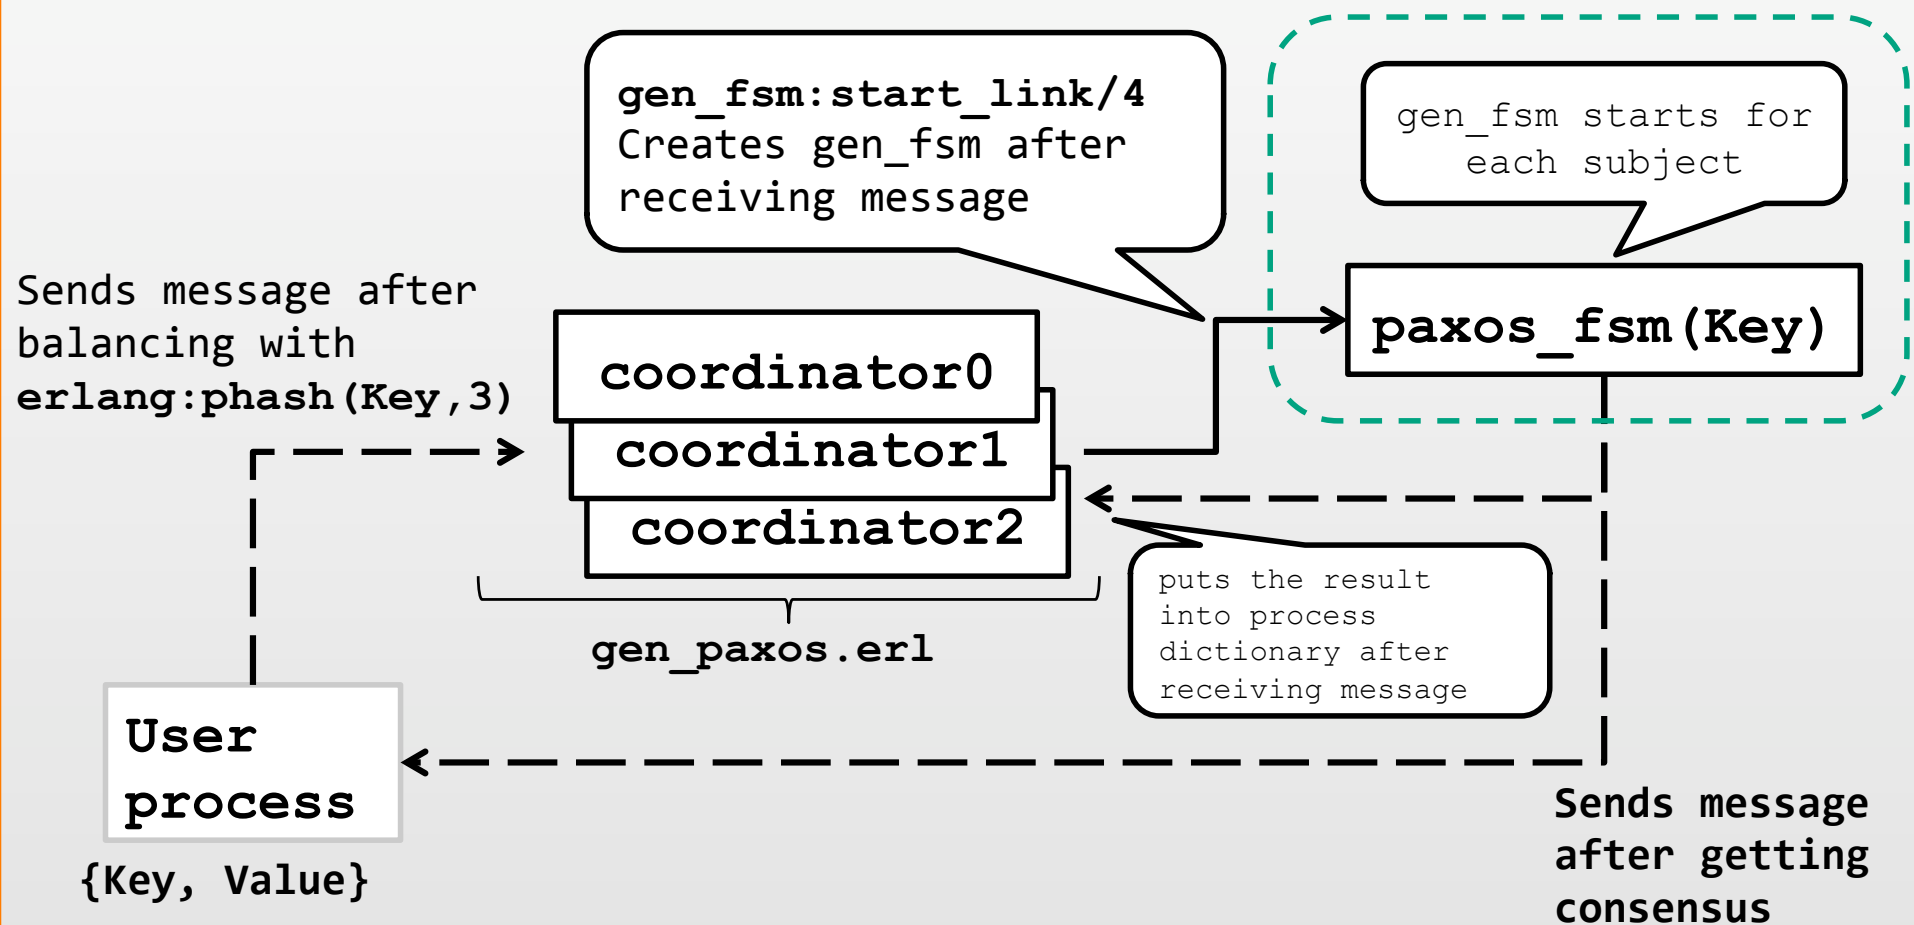
**Module:State(Event, Data)**  
-> {next\_state,  
NextState, NextData}

(See paxos\_fsm.erl for details)



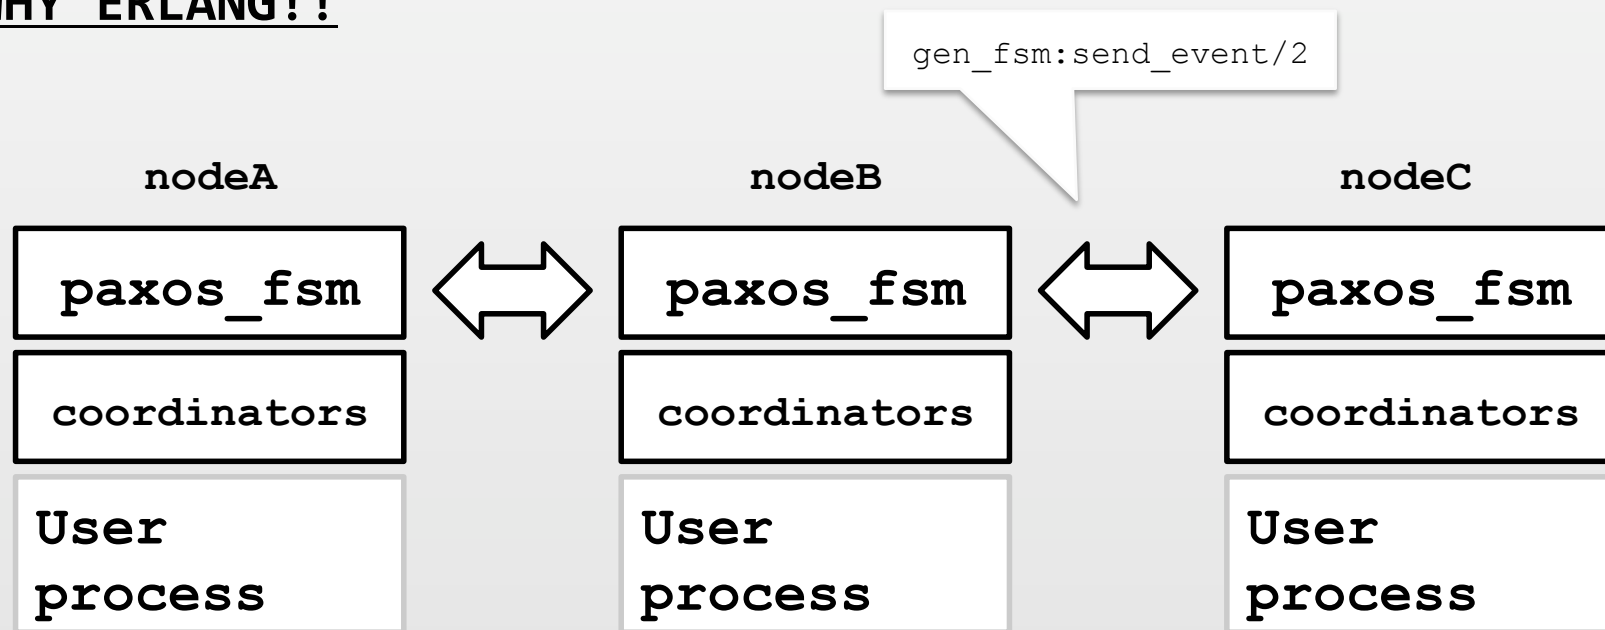
# Process Architecture of a Node (gen\_paxos.erl)

- A-little-bit load balancing
- Not pursuing consensus performance
- Thinkin'bout Chubby[5] (no dynamic node addition/removal)



## Messaging between nodes (paxos\_fsm.erl)

- FSMs talk each other by sending events with `globalname()`
  - `globalname() = {global, GlobalName}`
  - `GlobalName` is for first argument of `gen_fsm:start_link/4`
- `gen_fsm:send_event/2` sends Events to **remote nodes**
  - `gen_fsm:send_event({global, GlobalName}, Event).` - that's all.
- Anyway, no need to design application protocol, object serialization, nor GoF 'State, Observer' pattern! **That's WHY ERLANG!!**



# Sample code – paxos\_fsm.erl

```
acceptor( {prepare, {S, N, _V, From}}, StateData) when N < StateData#state.n -> %{{S, Nc, Vc,  
    send( From, S, { prepare_result, {S, StateData#state.n, StateData#state.value, node()}} ),  
    {next_state, acceptor, StateData, ?DEFAULT_TIMEOUT};  
  
%% acceptor( {prepare, {S, N, V, From}}, {{S, N, V}, Nums} ) -> % when N == Nc  
%%      {next_state, acceptor, {{S, N, V}, Nums}, ?DEFAULT_TIMEOUT};  
acceptor( {prepare, {S, N, V, From}}, StateData ) when N >= StateData#state.n ->  
    send( From, S, { prepare_result, {S, StateData#state.n, StateData#state.value, node()}} ),  
    {next_state, acceptor, StateData#state{n=N, value=V}, ?DEFAULT_TIMEOUT};
```

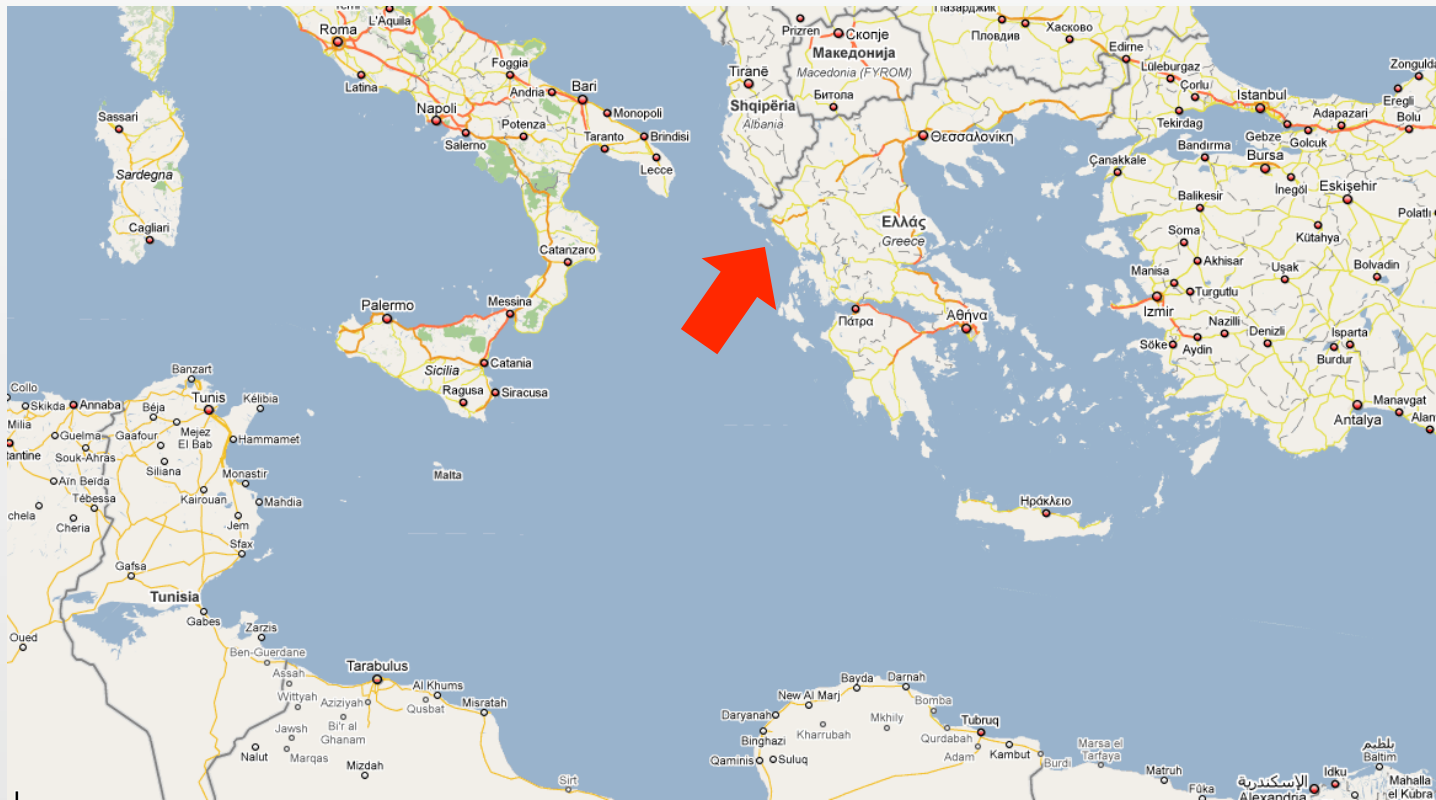
## 読み方

- acceptorのときに{prepare, {S, N, V, From}} (自分より大きなN)なイベントを受け取ったら、
- prepare\_resultを返信して
- NとVを更新してacceptor状態へ移る。

Where is PAXOS?

## Where is PAXOS?

- Paxos Isles in Aegean Sea[1]
  - ‘the part-time parliament’[3]



# Where is PAXOS? cont'd

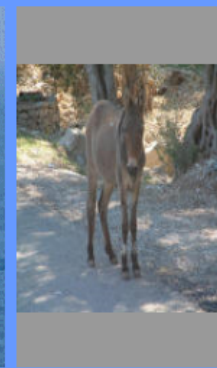
- Looks so nice place to visit!

[Paxos Weather](#)  
[Removals & Transport](#)  
[To Advertise or for](#)  
[Information on .....](#)  
[www.paxos-greece.com](http://www.paxos-greece.com)  
[Click here](#)

Just a few of the places and sites that you will find on this beautiful island of Paxos



*This web site is designed for the "Independent Traveller" and provides all the information you need to make your own holiday in Paxos. You may wish however, to leave arrangements to the Paxos Travel Agents listed above.*





## When?? - References

- [1]PAXOS island
  - <http://www.paxos-greece.com/>
- [2]Google, The Chubby Paper, 2006.
- [3]Leslie Lamport, 'The Part-Time Parliament', ACM Trans. Comput. Syst. Volume 16 , Issue 2 (May 1998) , pp. 133 - 169.
- [4]Leslie Lamport, 'Paxos Made Simple', SIGACT News, Vol. 32, No. 4. (December 2001), pp. 51-58.
- [5]T Chandra et al., 'Paxos made live: an engineering perspective', PODC '07, ACM Press, 2007, pp. 398-407.
- [6]Wikipedia, 'Paxos\_algorithm',
  - [http://en.wikipedia.org/wiki/Paxos\\_algorithm](http://en.wikipedia.org/wiki/Paxos_algorithm)
  - 1990年くらいにはもうネタはできていたらしい
- [7]Lamport's patent (さすがM\$...)
  - <http://www.j-tokkyo.com/2005/G06F/JP2005-196763.shtml>
- [8]David Mazieres, 'Paxos Made Practical', tech report, 2007.
  - <http://net.pku.edu.cn/~course/cs501/2008/reading.html>
- [9]Dwork, Cynthia et al., 'Consensus in the presence of partial synchrony', J. ACM 35-2, pp. 288-323, 1988.
- Other implementations
  - single-phase majority-voting is used in distributed memory, distributed CPU; rather hardware.
  - libpaxos

- Any Questions?



<http://creativecommons.org/licenses/by-sa/2.1/jp/>





# appendix

Known Bugs

FAQ

State Diagram

## Appendix – Known Bug(s)

- learnerの意味を勘違いした実装になっています
  - gen\_paxosでは、acceptした後になる状態のこと → 間違い
  - 本当は、議決には参加しないで結果に従うノードのこと
- <http://cooldaemon.tumblr.com/post/130422117/gen-paxos>
  - 幾つかのノードを起動し、幾つかの議題を出してみた。  
その後、過半数のノードを除去し、議題を出すと、ひたすら合意を取り続ける。  
この状態で、除去したノードと同じ名前のノードを復活させ、net\_adm\_ping/1 で繋ぐ。  
予想では、ここで合意が取れると思ったが、駄目だった。



## Appendix – FAQ

- 議題ごとに `gen_fsm` を使い捨てているのでどっかにプロセスをストックしておいて再利用できないもんなかなあ？（優先度低）
  - A. Erlangの思想から言って(プロセスは低コスト)、使い捨てすべきだと思います。今は作りっぱなしなので、どこかで破棄するタイミングがあるとよいかも。
- ノードの追加・除去に対応できるよう、ノードの一覧は外だしにした方が良くも。  
あとは、合意を取る最大回数も欲しい。
  - A. Chubbyはノードの動的追加・除去を前提としてないので外だしにはしていません。
  - A. `n`に上限を設けておいて、それをみんなが超えたら失敗、という作りはアリだと思います。
  - 結局、私に使い方(ユースケース)があまり見えていないままとりあえず書いたのが問題で、Chubby以外の便利な使い方があれば教えてください。
- 合意の間隔と回数は、起動時のグローバル値と、合意単位で設定するローカル値に分割すると使いやすいかも？
  - A. `timeout`時間はマクロ値(`?DEFAULT_TIMEOUT`)。早く決めたいときには短めに、確実に決めたいときは長めに設定できるように、`start_link`の引数とか`#state`のメンバにするのはよいかも。

# Appendix – State Diagram

ある議題についてのstate diagram

	states	$n' < n < n''$	返信	ブロードキャスト		
	nil(n,v) ignore	preparing (n, v) prepare( n, v )を送り返す	proposing (n,v) ignore	acceptor(n,v) prepare_result(n, v)を送る	learner(n,v) ignore	decided(n,v) ignore
prepare( n', v )						
prepare( n'', v )	prepare_result(0, nil)を返信してacceptor(n,v)へ	prepare_result(n, v)を送ってacceptor(n'', v'')へ	prepare_result(n, v)を送ってacceptor(n'', v'')へ	prepare_result(n, v)を送ってacceptor(n'', v'')へ	prepare_result(n,v)を送ってacceptor(n'', v'')へ	ignore
prepare_result(0, nil)	ignore	過半数を過ぎていたらpropose(n, v)を送ってproposing(n,v)へ	ignore	ignore	ignore	ignore
prepare_result( n', v )	ignore	過半数を過ぎていたらpropose(n, v)を送ってproposing(n,v)へ	ignore	ignore	ignore?	ignore
prepare_result( n, v )	ignore	過半数を過ぎていたらpropose(n, v)を送ってproposing(n,v)へ	ignore	-	-	ignore
prepare_result( n'', v'' )	ignore	prepare_result(n,v)を返信してacceptor(n'', v'')へ	prepare_result(n,v)を返信してacceptor(n'', v'')へ	ignore	ignore?	ignore
general_timeout	prepare(n,v)をみんなに送ってpreparing(n,v)へ	n++, nilへ	n++, nilへ	n++, nilへ	n++, nilへ	n++, nilへ
propose(n', v)	ignore	ignore	ignore	ignore	ignore	decide(n, v)を返信
propose(n, v)	ignore	-	-	propose_result(n, v)を返信してlearner(n, v)へ	-	ignore
propose(n'', v'')	ignore	propose_result(n, v)を返信してlearner(n'', v'')へ	propose_result(n, v)を返信してlearner(n'', v'')へ	propose_result(n, v)を返信してlearner(n'', v'')へ	propose_result(n, v)を返信してlearner(n'', v'')へ	ignore
propose_result(n', v')	ignore	ignore	ignore	ignore	ignore	reply decide(n,v)
propose_result(n, v)	ignore	ignore	過半数を超えれば全員にdecide(n,v)を送りend(n,v)へ	-	-	reply decide(n,v)
propose_result(n'', v'')	ignore	propose_result(n, v)を返信してlearner(n'', v'')へ	propose_result(n, v)を返信してlearner(n'', v'')へ	-	-	reply decide(n,v)
decide(n, v)	end(n,v)	end(n,v)	end(n,v)	end(n,v)	end(n,v)	end(n,v)
decide(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')	end(n'', v'')

- Thanks !



<http://creativecommons.org/licenses/by-sa/2.1/jp/>